

SLI: Sign Language Interpreter

Nour Hany¹, Laila Khaled², Yasmine El Qorashy³, and Ahmed Sameh⁴

¹College of Computer Science, Egypt University of Informatics, New Administrative Capital, 22-101068@students.eui.edu.eg

²College of Computer Science, Egypt University of Informatics, New Administrative Capital, 22-101078@students.eui.edu.eg

³College of Computer Science, Egypt University of Informatics, New Administrative Capital, 22-101174@students.eui.edu.eg

⁴College of Computer Science, Egypt University of Informatics, New Administrative Capital, 22-101198@students.eui.edu.eg

January 2025

Abstract

This project presents a comprehensive system for real-time Sign Language Translation at both the word and letter levels, with a focus on Egyptian Sign Language (ESL). For word-level SLI, we organized volunteers to collect a custom ESL video dataset and used MediaPipe to extract pose and hand landmarks from each frame. We then adopted a transfer-learning strategy: models were first pretrained on the large-scale WLASL dataset and subsequently fine-tuned on our collected ESL data. Several deep learning architectures were explored including single-stream and multistream Transformers, as well as LSTM models with and without attention and through systematic benchmarking we obtained a best model achieving 90% accuracy on word-level recognition.

At the letter level, we integrated YOLO-based hand detection with MediaPipe hand tracking to isolate fine-grained hand shapes, and trained deep learning classifiers to recognize individual letters. The resulting letter-level subsystem achieved 93% accuracy in controlled evaluation. Together, these components form an end-to-end, real-time Sign Language Interpreter capable of translating both isolated signs and fingerspelling, demonstrating the effectiveness of combining volunteer-collected ESL data, transfer learning from WLASL, and landmark-based deep sequence modeling for sign language understanding.

Keywords: Sign Language Interpreter, Deep Learning, Egyptian Sign Language, SLI, ESL, WASL, Transformer, LSTM, YOLO, Mediapipe

Contents

1	Introduction	1
1.1	Background	1
1.2	Problem Statement	1
1.3	Proposed Solution	2
1.4	Objectives	2
2	Letter-Level Sign Language Interpretation (SLI)	3
2.1	Methodology	3
2.1.1	Dataset	3
2.1.2	Data Preprocessing	4
2.1.3	Architecture	4
2.2	Results	6
2.2.1	Evaluation Metrics	6
2.2.2	Hyperparameters and Performance	7
2.2.3	Discussion	7
3	Word-Level Sign Language Interpretation (SLI)	10
3.1	Methodology	10
3.1.1	Dataset	10
3.1.2	Data Preprocessing	10
3.1.3	Architecture And Results	11
4	Conclusions and Future Work	17
4.1	Conclusion	17
4.2	Limitations	17
4.3	Future Work	18

List of Abbreviations

NN	Neural Network
LSTM	Long Short Term Memory
SLI	Sign Language Interpreter
GB	Extreme Gradient Boosting Regression Model
BiLSTM	Bidirectional Long Short-Term Memory
CNN	Convolutional Neural Network
ASL	American Sign Language
ESL	Egyptian Sign Language

Chapter 1

Introduction

1.1 Background

Communication is central to human interaction, enabling social integration, access to information, and meaningful participation in daily life. For millions of deaf and hard-of-hearing individuals worldwide, sign language is the primary medium through which this communication occurs. Unlike spoken language, sign languages encode meaning through complex spatiotemporal patterns involving hand shapes, motion trajectories, facial expressions, and body posture. Effective computational interpretation therefore requires models capable of understanding both fine-grained spatial features and temporal dynamics.

The global need for accessible communication technologies continues to grow. According to the World Health Organization, more than 430 million people currently experience disabling hearing loss, a number expected to exceed 700 million by 2050. Despite this, automated tools for sign language interpretation remain limited, especially for low-resource regional dialects such as Egyptian Sign Language (ESL), which has no publicly available dataset and differs significantly from widely studied languages like American Sign Language (ASL). This gap limits participation in education, employment, and essential services for members of the Deaf community.

1.2 Problem Statement

Although sign language is a rich and expressive linguistic system, communication between deaf individuals and the hearing population remains challenging due to the scarcity of reliable, real-time interpretation tools. Human interpreters are not always accessible, and text-based alternatives may fail to capture nuances of meaning or context. Existing automated translation systems are predominantly designed for ASL and perform poorly when applied to other dialects, especially those with different gesture vocabularies or stylistic variations.

The lack of scalable, dialect-specific, and real-time sign language translation solutions contributes to persistent communication barriers. These challenges are amplified for low-resource languages like ESL, where the absence of annotated datasets and pretrained models makes traditional supervised learning approaches infeasible. This motivates the need for approaches that combine transfer learning, landmark-based feature extraction, and deep sequence modeling to enable accurate and efficient translation under data scarcity.

1.3 Proposed Solution

This project adopts a two-stage deep learning approach to build a real-time sign language interpreter. The first stage targets letter-level recognition from static images of individual hand gestures, using computer vision models to classify fine-grained spatial features.

The second stage extends the system to word-level interpretation from video sequences, introducing sequence models to capture temporal dynamics in continuous signs. Because no public datasets exist for Egyptian Sign Language (ESL), we organized volunteers to collect a custom ESL dataset and applied transfer learning: models are first trained on WLASL and then fine-tuned on our collected ESL data.

1.4 Objectives

The primary objective of this project is to develop an automated system that translates Egyptian Sign Language (ESL) into text in real time, helping bridge the communication gap between deaf or hard-of-hearing individuals and the hearing population.

The specific objectives of this project are as follows:

- **Data Collection and Transfer Learning:** Build a custom ESL dataset through volunteer recordings and leverage transfer learning by pretraining models on the WLASL dataset and fine-tuning them on the collected ESL data.
- **Model Exploration and Optimization:** Evaluate and compare multiple deep learning architectures, including LSTM networks (with and without attention) and Transformer-based models, in both single-stream and multistream variants, to identify the most effective design for word-level recognition.
- **Landmark-Based Feature Extraction:** Use MediaPipe to extract robust hand and body pose landmarks, transforming raw video into structured spatio-temporal features suitable for sequence modeling.
- **Real-Time Prediction and Sliding-Window Inference:** Develop a real-time prediction pipeline that processes live video using sliding windows, outputs text as soon as confidence thresholds are met, and supports continuous signing.
- **Gesture Detection and End-of-Sign Handling:** Incorporate mechanisms to detect the presence or absence of gestures, automatically infer when a sign has ended (e.g., hands lowered), and pad or trim sequences to maintain temporal consistency.
- **Practical Deployment:** Design a lightweight, efficient implementation suitable for deployment on standard consumer hardware (e.g., laptops), ensuring the system is accessible and usable in real-world settings.

By achieving these objectives, the project aims to advance assistive technologies for the Deaf community and demonstrate how transfer learning and landmark-based deep sequence models can enable real-time interpretation in low-resource sign languages such as ESL.

Chapter 2

Letter-Level Sign Language Interpretation (SLI)

2.1 Methodology

2.1.1 Dataset

The dataset used for letter-level sign language interpretation is a combination of merged datasets sourced from RoboFlow. This dataset comprises static images of hand gestures representing the 26 letters of the English alphabet. Each image corresponds to a specific letter class, as depicted in Figure [2.2](#).

The total dataset consists of:

- **2,268 images**
- **Total size:** 45.16 MB
- **Classes:** 26 (A-Z)



Figure 2.1: Collage of hand gestures representing the English alphabet (A-Z).

The images capture different hand configurations and angles to ensure diversity in training data. The dataset covers various lighting conditions and backgrounds to enhance model generalization.

2.1.2 Data Preprocessing

The dataset underwent minimal preprocessing, as the images were already resized to a uniform frame size during the data collection process. This ensured consistency across all samples, eliminating the need for additional resizing or cropping steps.

The only additional preprocessing step applied to the images was pixel normalization. Each pixel value was scaled to the range $[0, 1]$ by dividing by 255. This normalization ensures that all features contribute equally to the model during training, facilitating faster convergence and improving overall performance.

2.1.3 Architecture

The system architecture for letter-level SLI consists of two main stages:

- **Feature Extraction**
- **Classification Model**

To implement this architecture, two distinct approaches were utilized: **YOLOv8** and **MediaPipe + Feedforward Neural Network (FNN)**. Each approach leverages different tools and methodologies to extract features and classify hand gestures.

Using YOLO (You Only Look Once)

YOLO is a state-of-the-art object detection algorithm designed for real-time performance. Initially introduced in 2015, YOLO revolutionized object detection by treating it as a regression problem rather than a classification task. This algorithm divides an image into grids, predicts bounding boxes, and assigns class probabilities using a single convolutional neural network (CNN).

YOLO Architecture Breakdown:

- **Residual Blocks:**
 - The image is divided into an $N \times N$ grid of cells.
 - Each grid cell is responsible for detecting objects that appear within its boundaries.
- **Bounding Box Regression:**
 - YOLO predicts bounding box coordinates (x , y , width, height) and assigns class probabilities.
- **Intersection Over Union (IoU):**
 - IoU evaluates the overlap between predicted and ground truth boxes, ensuring only the most accurate detections are retained.
- **Non-Maximum Suppression (NMS):**
 - When multiple bounding boxes detect the same object, NMS retains the box with the highest confidence score while discarding overlapping boxes.

These processes allow YOLO to achieve high-speed, real-time hand detection and gesture classification.

Using MediaPipe + Feedforward Neural Network (FNN)

Feature Extraction: MediaPipe's Hand Tracking module detects 21 hand landmarks (knuckle and finger points). These landmarks are represented as 2D coordinates (x , y) and converted into a 42-dimensional feature vector. This compact vector captures the spatial relationships between key points, which are crucial for gesture recognition.

Classification Model (FNN): A Feedforward Neural Network (FNN) is used to classify hand gestures based on the extracted features. The FNN processes the 42-dimensional input vector through:

- Two dense layers, each activated by ReLU (Rectified Linear Unit).
- The final output layer is the softmax layer contains 26 neurons (A-Z), representing each letter of the alphabet.
- The softmax activation function assigns probabilities to each class.

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	11,008
dense_1 (Dense)	(None, 128)	32,896
dense_2 (Dense)	(None, 26)	3,354

Total params: 47,258 (184.60 KB)

Trainable params: 47,258 (184.60 KB)

Non-trainable params: 0 (0.00 B)

Figure 2.2: Mediapipe + FNN model architecture.

This architecture efficiently classifies gestures while maintaining low computational overhead.

2.2 Results

2.2.1 Evaluation Metrics

To evaluate and compare the performance of the YOLO and Mediapipe-FNN architectures, the following metrics were employed:

1. **Accuracy:** Measures the overall percentage of correctly classified letters.

$$Accuracy = \frac{\text{Correctly Recognized Samples}}{\text{Total Samples}} \times 100 \quad (2.1)$$

2. **F1-Score:** Balances precision and recall, providing a robust assessment of classification performance.

$$Precision = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}} \quad (2.2)$$

$$Recall = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}} \quad (2.3)$$

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (2.4)$$

Figure 2.5 shows the accuracy and loss curves for the Mediapipe-FNN model during training, while **Figure 2.6** presents the confusion matrix, highlighting the model's classification performance across all letter classes.

2.2.2 Hyperparameters and Performance

The following table summarizes the hyperparameters and performance metrics for both the YOLO and Mediapipe-FNN models:

Model	Optimizer	Epochs	Batch Size	Loss Function	Test Accuracy
YOLOv8	ADAM	200	16	Categorical Cross Entropy	69.8%
Mediapipe + FNN	ADAM	200	16	Categorical Cross Entropy	93.33%

Table 2.1: Hyperparameters and Performance Metrics for YOLOv8 and Mediapipe-FNN Models.

2.2.3 Discussion

The results demonstrate that:

- Mediapipe-FNN significantly outperforms YOLOv8 in terms of test accuracy (93.33
- Mediapipe's ability to extract precise hand landmarks directly contributes to higher accuracy in gesture classification.




Figure 2.3: Accuracy and Loss curves for Mediapipe-FNN during training.




Figure 2.4: Confusion Matrix for Mediapipe-FNN (Letter Classification).




Figure 2.5: Accuracy and Loss curves for YOLO during training.




Figure 2.6: Confusion Matrix for YOLO (Letter Classification).

Chapter 3

Word-Level Sign Language Interpretation (SLI)

3.1 Methodology

3.1.1 Dataset

This project relies on a combination of publicly available data and a custom-collected dataset tailored for Egyptian Sign Language (ESL). To develop and evaluate the model architecture, we first utilized a processed version of the WLASL (Word-Level American Sign Language) dataset from Kaggle, which contains approximately 12,000 video samples across 2,000 sign classes. Each class includes multiple short video clips recorded by different signers, providing useful variability for pretraining. The dataset is available at: [Kaggle – WLASL Processed](#).

While WLASL offers a valuable large-scale corpus for transfer learning, no public datasets exist for Egyptian Sign Language. To address this gap, we organized volunteer participants to record ESL signs using consistent guidelines for lighting, framing, and gesture clarity. These recordings were manually verified and annotated, forming a custom ESL dataset used for fine-tuning. This two-stage data strategy of pretraining on WLASL and fine-tuning on volunteer-collected ESL data enables effective model performance in a low-resource setting while ensuring the final system is tailored to real ESL gestures.

3.1.2 Data Preprocessing

Video-to-frame conversion: On average, the videos contained approximately 69 frames. To standardize the data, all videos were uniformly reduced to 40 frames. For videos exceeding 40 frames, frames were sampled evenly across the video. For example, if a video had 120 frames, every 3rd frame was selected to ensure balanced coverage. Additionally, the first and last five frames were removed to eliminate motionless segments. For videos with fewer than 40 frames, padding was applied by repeating the last frame to ensure consistency across the dataset.




Figure 3.1: Video Frames Distribution.

Frame resizing: All frames were resized to 480x320 pixels while preserving their aspect ratio to maintain the integrity of the hand landmarks. Missing pixels resulting from resizing were padded with blank to create uniform frame dimensions.

Augmentation: To mitigate overfitting, data augmentation techniques such as random rotations, random zooming, shear transformations, and horizontal flipping were applied to the each set of frames.

3.1.3 Architecture And Results

Sign language interpretation at the word level involves translating a video into a word. Unlike letter-level interpretation, which deals with static images, word-level interpretation processes a sequence of frames forming a video. This temporal aspect requires sequential models capable of learning and interpreting time-series data. To handle this, the architecture integrates models designed to process frame sequences efficiently.

We experimented with several models, including:

- **Single-stream Transformer (Custom)**
- **Multistream Transformer (Custom)**
- **LSTM with Attention**
- **LSTM without Attention**

The following subsections provide an in-depth explanation of each architecture, accompanied by diagrams depicting the model structure.

Single-stream Transformer

The single-stream transformer processes input from a single stream, where the pose, left-hand, and right-hand features are concatenated before being passed into the transformer layers. This method leverages the self-attention mechanism to capture temporal dependencies across frames.

While the transformer architecture demonstrated the ability to learn from the training data effectively, the model exhibited significant overfitting. This is evident from the disparity between the training and validation performance metrics, as shown in Figure 3.3.

In the left plot of Figure 3.3, the training loss steadily decreases and eventually reaches near zero, indicating that the model fits the training data exceptionally well. However, the validation loss fluctuates and begins to rise after approximately 100 epochs, suggesting that the model fails to generalize to unseen data. The right plot further illustrates this overfitting, where the training accuracy approaches 100%, while the validation accuracy plateaus around 65-70%.

This performance gap indicates that the model memorized the training data rather than learning generalizable patterns. Increasing the dataset size, applying stronger regularization techniques, and employing data augmentation could help mitigate this overfitting issue.




Figure 3.2: Training and Validation Loss/Accuracy for the Single-stream Transformer.




Figure 3.3: Training and Validation Loss/Accuracy for the Single-stream Transformer.

Multistream Transformer

In the multistream transformer, pose, left-hand, and right-hand features are processed in separate streams through independent transformer blocks. The outputs are concatenated before passing through the dense layers for classification. This architecture aims to preserve the independent importance of each stream while allowing cross-stream attention.

Despite the architectural advantages, the multistream transformer performed poorly, similar to the single-stream approach. As depicted in Figure 3.4, the model exhibits severe overfitting, characterized by a large disparity between the training and validation performance.

The left plot shows that the training loss consistently decreases to near zero, while the validation loss fluctuates significantly and increases steadily after approximately 100 epochs. This indicates that the model memorized the training data but failed to generalize to new, unseen data. On the right plot, the training accuracy reaches 100%, but the validation accuracy stagnates around 50-55%, highlighting the inability of the model to adapt to the validation set.

The poor generalization is likely due to the relatively small dataset size, as transformers typically require large-scale data to perform effectively. Additionally, the complexity introduced by handling three separate streams might have contributed to the unstable validation results.




Figure 3.4: Training and Validation Loss/Accuracy for the Multistream Transformer. Severe overfitting is evident as the validation performance deteriorates while training accuracy approaches 100%.

LSTM with Attention

The LSTM with attention architecture employs a bidirectional LSTM to process the frame sequences, followed by an attention mechanism that focuses on the most critical frames. This allows the model to emphasize key temporal patterns, enhancing accuracy by selectively attending to important segments of the input sequence.

The LSTM with attention model achieved a test accuracy of **85%**, outperforming both the single-stream and multistream transformer models. This performance proves that simpler models usually generalize better with a relatively small dataset.

Layer (type)	Output Shape	Param #	Connected to
input_layer_6 (InputLayer)	(None, 30, 33, 3)	0	-
input_layer_7 (InputLayer)	(None, 30, 21, 3)	0	-
input_layer_8 (InputLayer)	(None, 30, 21, 3)	0	-
time_distributed_3 (TimeDistributed)	(None, 30, 99)	0	input_layer_6[0]_
time_distributed_4 (TimeDistributed)	(None, 30, 63)	0	input_layer_7[0]_
time_distributed_5 (TimeDistributed)	(None, 30, 63)	0	input_layer_8[0]_
concatenate_1 (Concatenate)	(None, 30, 225)	0	time_distributed_ time_distributed_ time_distributed_
bidirectional (Bidirectional)	(None, 30, 256)	362,496	concatenate_1[0]_
batch_normalization (BatchNormalizatio..)	(None, 30, 256)	1,024	bidirectional[0]_
dropout_10 (Dropout)	(None, 30, 256)	0	batch_normalizat_
lstm_1 (LSTM)	(None, 30, 64)	82,176	dropout_10[0][0]
dense_9 (Dense)	(None, 30, 128)	8,320	lstm_1[0][0]
dropout_11 (Dropout)	(None, 30, 128)	0	dense_9[0][0]
dense_10 (Dense)	(None, 30, 4)	516	dropout_11[0][0]

** Total params: 454,532 (1.73 MB)
 ** Trainable params: 454,020 (1.73 MB)
 ** Non-trainable params: 512 (2.00 KB)

Figure 3.5: LSTM with Attention Model Architecture.

LSTM without Attention

In this architecture, the model employs a bidirectional LSTM without an attention layer. The frame sequences are processed sequentially, and the final hidden state is used for classification. This simpler approach, although lacking the attention mechanism, surprisingly outperformed the LSTM with attention, achieving a test accuracy of **90%**.

The superior performance of the LSTM without attention can be attributed to several factors:

- **Smaller Dataset:** Attention mechanisms generally require larger datasets to learn effectively. With limited training data, the attention layer may introduce unnecessary complexity, leading to overfitting.
- **Reduced Noise Amplification:** Attention mechanisms may amplify minor variations or

noise in the data, especially in smaller datasets. Without attention, the LSTM processes frames uniformly, which may prevent the model from overfitting to unimportant details.

Despite its strong performance, the model could potentially benefit from attention mechanisms if more data were available. This highlights the importance of dataset size in determining the suitability of advanced architectures.




Figure 3.6: LSTM without Attention Model Architecture. Training and validation curves show minimal overfitting, with validation accuracy stabilizing around 90%.

Layer (type)	Output Shape	Param #	Connected to
input_layer_9 (InputLayer)	(None, 30, 33, 3)	0	-
input_layer_10 (InputLayer)	(None, 30, 21, 3)	0	-
input_layer_11 (InputLayer)	(None, 30, 21, 3)	0	-
time_distributed_6 (TimeDistributed)	(None, 30, 99)	0	input_layer_9[0]-
time_distributed_7 (TimeDistributed)	(None, 30, 63)	0	input_layer_10[0]-
time_distributed_8 (TimeDistributed)	(None, 30, 63)	0	input_layer_11[0]-
concatenate_2 (Concatenate)	(None, 30, 225)	0	time_distributed_ time_distributed_ time_distributed_
bidirectional_1 (Bidirectional)	(None, 30, 256)	362,496	concatenate_2[0]-
batch_normalizatio_ (BatchNormalizatio_)	(None, 30, 256)	1,024	bidirectional_1[-
dropout_12 (Dropout)	(None, 30, 256)	0	batch_normalizat_
lstm_3 (LSTM)	(None, 64)	82,176	dropout_12[0][0]
dense_11 (Dense)	(None, 128)	8,320	lstm_3[0][0]
dropout_13 (Dropout)	(None, 128)	0	dense_11[0][0]
dense_12 (Dense)	(None, 4)	516	dropout_13[0][0]

Total params: 454,532 (1.73 MB)

Trainable params: 454,020 (1.73 MB)

Non-trainable params: 512 (2.00 KB)

Figure 3.7: LSTM without Attention Architecture.

Chapter 4

Conclusions and Future Work

4.1 Conclusion

This study highlights the transformative potential of AI for advancing word-level sign language interpretation, with a focus on Egyptian Sign Language. The proposed system translates video sequences into words by leveraging **MediaPipe** for robust pose and hand landmark extraction, **LSTM**-based sequence models to capture temporal dynamics, and a **feedforward neural network** for final classification. This pipeline effectively models the spatio-temporal nature of sign gestures.

The letter-level implementation served as a foundational stage, enabling reliable preprocessing and feature extraction that supported the more complex word-level task. By combining transfer learning from WLASL with a volunteer-collected ESL dataset, the project demonstrates that accurate recognition is achievable even in low-resource settings. Despite constraints related to dataset size, diversity, and hardware limitations, the system establishes a practical framework for scalable, real-time sign language translation.

While further work is needed to expand the dataset, incorporate richer multimodal cues (such as facial expressions and full-body movement), and optimize deployment for a wider range of devices, the results presented here represent a meaningful step toward bridging communication gaps for **deaf and hard-of-hearing individuals**. This integration of state-of-the-art tools and community-driven data collection underscores the role of AI as a powerful enabler of accessible and inclusive communication technologies.

4.2 Limitations

1. **Dataset Quality:** The dataset's limited size and diversity, particularly in the number of videos per word, constrained the model's ability to generalize effectively. This was especially challenging for sequential models like transformers, which require larger datasets.
2. **Hardware Constraints:** The training process was limited by the computational capabilities of the hardware used, affecting both speed and the ability to experiment with more complex architectures or larger datasets.
3. **Real-Time Deployment:** For real-life applications, the system's inference speed must be improved to ensure efficient real-time sign language interpretation.

4.3 Future Work

1. **Dataset Expansion and Community Contributions:** As Egyptian Sign Language lacks publicly available datasets, future work will focus on expanding the volunteer-collected corpus. Researchers, practitioners, and members of the Deaf community who are interested in contributing recordings or collaborating on dataset development are encouraged to reach out to the project team at esl.dataset@gmail.com. Expanding the dataset is essential for improving model robustness and supporting broader linguistic coverage.
2. **Multimodal Translation:** Incorporate facial expressions and body movements to enhance translation accuracy and context understanding. This multimodal approach can capture the full complexity of sign language.
3. **Language Expansion:** Extend the system to support multiple languages, enabling broader accessibility and impact. Including more underrepresented sign languages could significantly increase the system's reach.
4. **System Deployment and Integration:** Focus on deploying and integrating this system into platforms and environments where it can be actively used.

References

- A. H. Althubiti and H. Algethami. Dynamic gesture recognition using a transformer and mediapipe. *International Journal of Advanced Computer Science and Applications*, 2024.
- T. Ananthanarayana et al. Deep learning methods for sign language translation. *ACM Transactions on Accessible Computing*, 2021.
- Mediapipe Documentation. *Real-Time Hand Tracking and Gesture Recognition Framework*. Google Developers, 2023.
- A. Hussain et al. An efficient and robust hand gesture recognition system. *Computer Science and Engineering*, 2023.
- A. Vaswani et al. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.
- World Health Organization. Deafness and hearing loss, 2023. URL <https://www.who.int/news-room/fact-sheets/detail/deafness-and-hearing-loss#:~:text=Overview,will%20have%20disabling%20hearing%20loss>. Accessed: Jan. 4, 2025.
- M. Zakariah et al. Sign language recognition for arabic alphabets using transfer learning. *Computational Intelligence and Neuroscience*, 2022.